

SEGMENT 08 • THE HONEST MAP

What you can build. What you cannot. Yet.

Twenty minutes of honesty buys more trust than twenty minutes of enthusiasm

Capabilities • limitations • the gotchas list • what happens next

Dr. Bhaveshkumar C. Dharmani • PhD (ICT), DA-IICT Gandhinagar • Founder, AI Vidhya4Sarvam

<https://www.aividhya.in/> • <https://www.aividhya4sarvam.in/> • bhavesh@aividhya.in • <https://www.linkedin.com/in/bhaveshdharmani/>

Genuinely strong today



Indic speech, both directions

23 languages in, 11 out. Code-mixed input handled natively. Nothing else comes close on Indian language coverage.



8 kHz telephony audio

Real Indian call traffic. Most Western models degrade badly here; this stack was built for it.



Indic document extraction

Typed, scanned and handwritten, across 23 languages and scripts, at ₹0.50 a page.



Cost

The Indic-native flagship is also the cheapest chat model on the platform. Document processing is 3x cheaper than a year ago.



Sovereign deployment

Managed, VPC, on-premise, air-gapped. SOC 2 Type II, ISO 27001, DPDP. Chanakya for defence-grade isolation.



Open weights

30B and 105B under Apache 2.0. You are not locked in — you can take the weights and run.



For any product where Indian language, Indian data residency, or Indian cost structure is a hard requirement, this is now the default answer rather than the patriotic one.

Where it is genuinely weaker

1

Frontier general reasoning still trails the best closed models

Sarvam-105B is excellent and competitive in its class. For the hardest general reasoning tasks in English, the global frontier is still ahead. Pretending otherwise costs you credibility.

2

Ecosystem depth is younger

Fewer Stack Overflow answers, fewer blog posts, fewer people who have hit your exact bug. You will be reading source and docs more often.

3

Documentation lags announcements

V4 and Vision 2.0 shipped ten days ago and are not in the public API docs. You saw this yourself in segment two.

4

TTS covers 11 languages, ASR covers 23

You can transcribe Santali. You cannot speak it back yet. Design around the narrowest model in your chain.

5

Hard operational limits

Document AI: 10 pages a job, 10 requests a minute on every plan. STT REST: 30 seconds. These shape your architecture, not just your code.



Say all of this out loud when you teach it. A room that hears you name real weaknesses believes you about the strengths.

The gotchas list

Gotcha	Symptom	Fix
client.chat.completions.create(...)	AttributeError	client.chat.completions(...) — no .create
reasoning eats max_tokens	content is None	Raise max_tokens or reasoning_effort=None
output_script on sarvam-translate	200 OK, wrong script	Use Mayura, or post-process
language_code on Document AI	Silently ignored	Use language=
output_format="markdown"	400 Bad Request	Use "md"
schema as a dict	'dict' has no attribute 'read'	json.dumps(schema)
file=open(...) on doc_ai	Type error	file=[(name, handle, mime)] — an array
sample_rate= on REST transcribe()	TypeError: unexpected keyword	Drop it — the WAV header carries it
sample_rate MISSING when streaming	Garbage transcript, no error	connect(sample_rate=...) — raw PCM has no header
Only checking status == completed	Pages silently lost	Handle partially_completed
Results before terminal state	409	Poll get_status first
Bulbul v3 by default	2x the TTS bill	v2 is often adequate for IVR
Assuming model defaults are stable	Behaviour changes silently	Pass model= explicitly, always
async with on the sync SarvamAI	asynchronous context manager	Use with — or AsyncSarvamAI
MSD+ and all other non-TTS text	Bill goes to ₹2000 a month	WAV with text-to-speech is more consistent



Screenshot this one. It is the single most useful artefact from today, and every item on it cost somebody a working day to discover.

"Why not just use OpenAI or Gemini?"



Where Sarvam genuinely wins

- Indic speech recognition — 23 languages, code-mixed, 8kHz telephony
- Indic TTS quality and voice range
- Document extraction across Indian scripts, including handwriting
- Data residency, DPDP, and India-domiciled processing
- Rupee pricing, GST invoices, local procurement
- On-premise and air-gapped deployment options
- Cost — meaningfully cheaper for these workloads



Where the global stack is still ahead

- Frontier general reasoning in English
- Multimodal breadth — video, advanced vision, real-time interruption
- Ecosystem maturity: tooling, integrations, community answers
- Very long context handling at the extreme end
- Model variety — you cannot A/B five frontier models here
- Documentation completeness and stability



The right answer is not loyalty to one stack. It is knowing precisely which workload belongs where — and being able to defend that choice to a risk team.

The other questions you will get

1

"Will Sarvam compete with what I build?"

They sell platforms to large enterprises with forward-deployed teams. Vertical depth, SMB, workflow-adjacent and on-prem appliance remain genuinely open. Be straight that this could change — and architect for portability.

2

"What is the real accuracy on my dialect?"

I do not know, and nor does anyone who has not tested it. Take twenty recordings from your own users and measure it in an afternoon. That answer beats any benchmark.

3

"What if they raise prices?"

30B and 105B are Apache 2.0. The API is OpenAI-compatible. Both facts mean real portability. Architect for model-swappability from day one and this stops being a risk.

4

"How long to production?"

Weeks for a prototype. Months for a regulated deployment. In Indian enterprise the long pole is procurement and security review, not engineering.

★

"I don't know, and here is how we would find out" is a better answer than a confident guess. Experienced rooms can tell the difference instantly.

Nobody asks a carpenter why he owns two saws

A carpenter has a hand saw and a circular saw. Nobody asks him to justify the choice, or accuses him of disloyalty to one.

He uses the circular saw for long straight cuts in cheap timber, and the hand saw when he is working close to a finished edge and cannot afford a mistake.

The question is never "which saw is better". It is "what am I cutting".



You will use Sarvam for Indic speech, Indian documents and anything a compliance officer has to sign off. You will use something else for frontier English reasoning. Knowing which is which is the skill — not picking a side.



What happens next

India's AI Sovereignty Month — 15 August to 15 September

Your next 30 days

1

Week 1

Ship something. Show one person in your ICP. Write down exactly what they said.

2

Week 2

Build a 10-case eval set from real data. Measure accuracy on YOUR audio.

3

Week 3

Price it. Send one paid-pilot proposal to one named prospect.

4

Week 4

Apply to the Sarvam Startup Program. Post your build in Discord. Contribute a cookbook example.



The one that compounds

Contributing to the cookbook costs you an evening and permanently changes how the ecosystem sees you. It is the cheapest developer marketing available, and it puts your name in front of the team building this.



Set up the tooling tonight — Agent Skills, MCP server, Kivi. Ten minutes now saves an afternoon next week.

India's AI Sovereignty Month

1

17-28 August — Campus tour

Free hands-on sessions for students at colleges across the region. If you are at an institution that would want this, message me tonight. I am not charging students.

2

29-30 August — The full workshop

Two days, 16 hours, ₹1,000. You build a voice agent that answers a real phone call, a document pipeline, an agentic workflow with evals, and your own costed business model.

3

19 September — 8-week cohort

From 'I can call the API' to 'I have shipped a product'. Weekly live sessions, code review, capstone. ₹2,500 students / ₹5,000 professionals. Early bird open now.

★

The ₹1,000 from the August workshop is credited against your cohort fee. If today was useful, that is the path — and the links are in chat now.

?


One ask before you go

If you are at a college, or know a coding club that would want a session like this — message me. My room is my pipeline, and I would rather ask than assume.

Everything you need, in one place

What	Where
Documentation	docs.sarvam.ai — append .md to any page for Markdown
Full docs as one file	docs.sarvam.ai/lms-full.txt
Dashboard, keys, playground	dashboard.sarvam.ai
Changelog — check before every project	docs.sarvam.ai/api/getting-started/changelog
Cookbook	github.com/sarvamai/sarvam-ai-cookbook
Agent Skills	github.com/sarvamai/skills — <code>npx skills add sarvamai/skills</code>
MCP server	<code>uvx sarvam-mcp</code>
Indus	indus.sarvam.ai
Status page	status.sarvam.ai
Discord community	discord.com/invite/5rAsyktcs
Startup Program	sarvam.ai/startup-program
Sarvam Champions	sarvam.ai/sarvam-champions
Self-hosting on SageMaker	docs.sarvam.ai/api/self-hosted/introduction
Migration guides	docs.sarvam.ai/api/migrations — from ElevenLabs, Cartesia, Deepgram, Gemini

Before you close the tab



The one question I actually need answered honestly

"What is one thing I should change for the 29–30 August workshop?" I am running this five more times on campuses this month. Flattering feedback helps nobody — tell me what did not land.



And if you generated audio you liked, drop it in chat. I want to count how many Indian languages this room covered tonight.

“

The model is not the moat.
The workflow, the distribution
and the cost structure are.

Thank you, and happy Independence Day. — aividhya.in · dharmanicb@gmail.com

NEXT

29–30 August The full workshop

16 hours. You build a voice agent, a document pipeline, an agentic workflow, and a business model.

₹1,000 — credited against the September cohort • Links in chat • aividhya.in